

Movie_Review_GMM

March 3, 2022

1 Unsupervised Sentiment Analysis using GMM

Extracting Tf-Idf feature vector from the given movie review dataset and Using GMM and EM algorithm to perform Unsupervised Sentiment Analysis

Libraries used: 1) Numpy - for numerical computations such as `fft()`, `dot` operator 2) Scipy - to find the likelihood of data points 3) Matplotlib - to plot the spectrogram

```
[1]: import numpy as np
import string, re, random
import matplotlib.pyplot as plt
from scipy.stats import multivariate_normal
```

1.0.1 Function `preprocess_text()`

Gets the vocabulary of words from the given text dataset

Strips the sentences into individual words

```
[2]: vocabulary = []
target = []
words_count = []
reviews = []

def preprocess_text(filename):
    split_by_char =[:, '/', '']

    with open(filename, "r") as f:
        line_number = 0
        for line in f.readlines():
            review = []
            words = line.split(" ")

            word_count = 0
            for word in words[:-1]:
                if word == "" or word == " " or word in string.punctuation:
                    continue

                word = word.strip(")(.*-&?:;, ")
```

```

        for char in split_by_char:
            if char in word:
                word2 = word.split(char)
                for word in word2:
                    if word.lower() not in vocabulary:
                        vocabulary.append(word.lower())
                    word_count += 1
                review.append(word.lower())
            else:
                if word.lower() not in vocabulary:
                    vocabulary.append(word.lower())
                review.append(word.lower())
                word_count += 1
        line_number += 1

    words_count.append(word_count)
    if "0" in words[-1]:
        target.append(0)
        review.append(0)
    else:
        target.append(1)
        review.append(1)
    reviews.append(review)

```

```

[3]: preprocess_text("movieReviews1000.txt")
     print("Number of unique words in vocabulary: ", len(vocabulary))

```

Number of unique words in vocabulary: 3165

Qs 4) a) Extracting Tf-Idf features per each word

1.0.2 Function calculate_idf()

Calculates the Inverse Document Frequency values for each word in the vocabulary

```

[4]: def calculate_idf(documents, vocabulary):

    idf = np.zeros((len(vocabulary), 1))

    n=len(documents)
    l = 0

    for word in vocabulary:
        word_count=0
        for words in documents:
            if word in words:
                word_count+=1

```

```

idf[l]=(np.log((1+n)/(word_count+1)))
l+=1
return idf

```

1.0.3 Function calculate_tf_idf()

Calculates the Term Frequency and Tf_Idf values for each word and document in the given text dataset

```

[5]: def calculate_tf_idf(idf):
    tf_idf = np.zeros((len(vocabulary), 0))

    line = 0
    for review in reviews:
        tf = np.zeros((len(vocabulary), 1))

        for word in review[:-1]:
            idx = vocabulary.index(word)
            tf[idx] += 1

        tf = (tf+1) / (2*len(review))
        line += 1
        tf_idf = np.hstack((tf_idf, tf))

    tf_idf_matrix = tf_idf * idf

    return tf_idf_matrix

```

```

[6]: idf = calculate_idf(reviews, vocabulary)
    tf_idf = calculate_tf_idf(idf)
    print("Size of the input feature matrix: ", tf_idf.shape)

```

Size of the input feature matrix: (3165, 1000)

Qs 4) b) Performing PCA To reduce the dimension of the given data

```

[7]: from sklearn.decomposition import PCA

    pca = PCA(n_components=10)

    dim_red_tf_idf = pca.fit_transform(tf_idf.T)

    print(dim_red_tf_idf.shape)

```

(1000, 10)

```

[8]: class KMeans():
    def __init__(self, n_clusters, max_iter, tol = 0.0001):
        self.n_clusters = n_clusters
        self.max_iter = max_iter
        self.tol = tol

        # X - shape (n_samples, n_features)
    def fit(self, X):

        self.centroid = [np.random.rand(X.shape[1], X.shape[1]) for i in
↪range(self.n_clusters)]

        data_points = random.sample(range(X.shape[0]), self.n_clusters)

        for i in range(self.n_clusters):
            self.centroid[i] = X[data_points[i]]

        epsilon = self.tol + 1
        for i in range(self.max_iter):

            if epsilon < self.tol:
                break

            self.clusters = [[] for i in range(self.n_clusters)]

            for sample in X:
                idx = np.argmin(np.array([np.linalg.norm(self.centroid[k] -
↪sample) for k in range(self.n_clusters)]))
                self.clusters[idx].append(sample)

            previous_centroid = self.centroid[:]

            epsilon = 0
            for j in range(self.n_clusters):

                if len(self.clusters[j]) == 0:
                    continue

                self.centroid[j] = np.mean(np.array(self.clusters[j]), axis = 0)

            err = abs(self.centroid[j] - previous_centroid[j])
            epsilon = max(epsilon, sum(err))

```

Initializing the mean and covariance for EM iteration using KMeans clustering

```

[9]: n_mixtures_ = 2
    covariance_type = 'diag'

```

```

clf = KMeans(n_clusters=2, max_iter = 100)

clf.fit(dim_red_tf_idf)

mu_k = [clf.centroid[0], clf.centroid[1]]
sigma_k = [[] for i in range(n_mixtures_)]

for i in range(n_mixtures_):
    sigma_k[i] = np.cov(np.array(clf.clusters[i]).T)

if covariance_type == 'diag':
    for i in range(n_mixtures_):
        diag = np.einsum('ii->i', sigma_k[i])
        save = diag.copy()
        sigma_k[i][...] = 0
        sigma_k[i] = np.diag(save)

```

Question 4) c) Training 2 mixture diagonal covariance GMM

```

[10]: class GMM:
    def __init__(self, n_mixtures, max_iter, covariance_type, tol = 0.0001):
        self.n_mixtures = n_mixtures
        self.max_iter = max_iter
        self.tol = tol
        self.covariance_type = covariance_type

    def initialize(self, mean, sigma):
        self.alpha = [(1/self.n_mixtures) for i in range(self.n_mixtures)]
        self.mean = mean
        self.sigma = sigma

    def expectation_step(self, X):
        self.likelihood_data = self.predict_likelihood(X)

    def maximization_step(self, X):
        self.alpha = np.mean(self.predict_likelihood(X), axis=0)

        for i in range(self.n_mixtures):
            like_prob = self.likelihood_data[:, i]
            num = X * like_prob.reshape((X.shape[0], 1))
            self.mean[i] = np.sum(num, axis = 0) / np.sum(like_prob)

            cov = like_prob.reshape((1, X.shape[0])) * (X - self.mean[i]).T
            cov = cov.dot((X - self.mean[i]))
            cov /= np.sum(like_prob)
            self.sigma[i] = cov

        if self.covariance_type == 'diag':

```

```

        for i in range(n_mixtures_):
            diag = np.einsum('ii->i', sigma_k[i])
            save = diag.copy()
            sigma_k[i][...] = 0
            sigma_k[i] = np.diag(save)

def predict_likelihood(self, X):
    likelihood = np.zeros( (X.shape[0], self.n_mixtures))

    for i in range(self.n_mixtures):
        distribution = multivariate_normal(mean=self.mean[i], cov=self.
→sigma[i])
        likelihood[:,i] = distribution.pdf(X)

    likelihood_gmm = likelihood * self.alpha
    total_likelihood = likelihood_gmm.sum(axis=1)[:, np.newaxis]
    likelihood_gmm = likelihood_gmm / total_likelihood
    return likelihood_gmm

def calc_log_likelihood(self, X):
    likelihood = np.zeros( (X.shape[0], self.n_mixtures) )
    for i in range(self.n_mixtures):
        distribution = multivariate_normal(
            mean=self.mean[i],
            cov=self.sigma[i])
        likelihood[:,i] = distribution.pdf(X)

    numerator = likelihood * self.alpha
    log_like = np.log(np.sum(numerator, axis=1))
    log_like = np.sum(log_like)
    return log_like

def plot_data(self, X):
    pred = self.predict(X)
    count = len([0 for i in pred if i == 1])
    print("Data points with maximum likelihood on Gaussian Mixture component_
→1: ", count)
    print("Data points with maximum likelihood on Gaussian Mixture component_
→2: ", 1000 -count)
    for i in range(X.shape[0]):
        if pred[i] == 1:
            plt.scatter(X[i][0], X[i][1], color='r')
        elif pred[i] == 0:
            plt.scatter(X[i][0], X[i][1], color='b')
    plt.xlabel("PCA Component 1")
    plt.ylabel("PCA Component 2")
    plt.show()

```

```

# X - n_samples, n_features
def fit(self, X):
    log_likelihood = []

    #last = self.max_iter

    for i in range(self.max_iter):
        self.expectation_step(X)
        self.maximization_step(X)

        print("Iteration: ", i+1)
        self.plot_data(X)

        log_likelihood.append(self.calc_log_likelihood(X))
        if i > 2 and abs(abs(log_likelihood[-1]) - abs(log_likelihood[-2])) <
→ self.tol:
            print("----- EM Iteration Converged -----")
            break

        #plt.scatter([i for i in range(last)], log_likelihood, marker="o",
→ color='g', linewidths=5)
        #plt.show()

    def predict(self, X):
        weights = self.predict_likelihood(X)
        return np.argmax(weights, axis=1)

```

```

[11]: gmm = GMM(n_mixtures = n_mixtures_, max_iter = 40, covariance_type =
→ covariance_type)
print(dim_red_tf_idf.shape)
gmm.initialize(mu_k, sigma_k)
gmm.fit(dim_red_tf_idf)


```

(1000, 10)

Iteration: 1

Data points with maximum likelihood on Gaussian Mixture component 1: 817


Data points with maximum likelihood on Gaussian Mixture component 2: 183



Iteration: 2

Data points with maximum likelihood on Gaussian Mixture component 1: 810


Data points with maximum likelihood on Gaussian Mixture component 2: 190



Iteration: 3

Data points with maximum likelihood on Gaussian Mixture component 1: 805


Data points with maximum likelihood on Gaussian Mixture component 2: 195



Iteration: 4

Data points with maximum likelihood on Gaussian Mixture component 1: 802


Data points with maximum likelihood on Gaussian Mixture component 2: 198



Iteration: 5

Data points with maximum likelihood on Gaussian Mixture component 1: 802


Data points with maximum likelihood on Gaussian Mixture component 2: 198



Iteration: 6

Data points with maximum likelihood on Gaussian Mixture component 1: 800


Data points with maximum likelihood on Gaussian Mixture component 2: 200



Iteration: 7

Data points with maximum likelihood on Gaussian Mixture component 1: 800


Data points with maximum likelihood on Gaussian Mixture component 2: 200



Iteration: 8

Data points with maximum likelihood on Gaussian Mixture component 1: 800


Data points with maximum likelihood on Gaussian Mixture component 2: 200



Iteration: 9

Data points with maximum likelihood on Gaussian Mixture component 1: 800


Data points with maximum likelihood on Gaussian Mixture component 2: 200



Iteration: 10

Data points with maximum likelihood on Gaussian Mixture component 1: 800


Data points with maximum likelihood on Gaussian Mixture component 2: 200



Iteration: 11

Data points with maximum likelihood on Gaussian Mixture component 1: 799


Data points with maximum likelihood on Gaussian Mixture component 2: 201



Iteration: 12

Data points with maximum likelihood on Gaussian Mixture component 1: 799


Data points with maximum likelihood on Gaussian Mixture component 2: 201



Iteration: 13

Data points with maximum likelihood on Gaussian Mixture component 1: 799


Data points with maximum likelihood on Gaussian Mixture component 2: 201



Iteration: 14

Data points with maximum likelihood on Gaussian Mixture component 1: 799

Data points with maximum likelihood on Gaussian Mixture component 2: 201



----- EM Iteration Converged -----

Question 4) d) Misprediction Rate

```
[12]: pred = gmm.predict(dim_red_tf_idf)

error = 0
for i in range(1000):
    if target[i] != pred[i]:
        error += 1
misclassification_percent = min(error, dim_red_tf_idf.shape[0] - error) /   
    ↪ dim_red_tf_idf.shape[0] * 100

print("Misclassification Percent: ", misclassification_percent, "%")
```

Misclassification Percent: 45.9 %

No, the true label does not correlates properly with the cluster identity

On average, there is a 45.9 % mismatch between true label and cluster identity